HHS Public Access

Author manuscript

Curr Opin Neurobiol. Author manuscript; available in PMC 2018 January 22.

Published in final edited form as:

Curr Opin Neurobiol. 2016 April; 37: 149–157. doi:10.1016/j.conb.2016.01.003.

Neural underpinnings of the evidence accumulator

Carlos D. Brody^{1,2} and Timothy D. Hanks^{3,4}

¹Howard Hughes Medical Institute

²Princeton Neuroscience Institute and Department of Molecular Biology, Princeton University, Princeton NJ 08540, USA

³Center for Neuroscience, University of California Davis, Davis, CA 95618, USA

⁴Department of Neurology, University of California Davis, Sacramento, CA 95817, USA

Abstract

Gradual accumulation of evidence favoring one or another choice is considered a core component of many different types of decisions, and has been the subject of many neurophysiological studies in non-human primates. But its neural circuit mechanisms remain mysterious. Investigating it in rodents has recently become possible, facilitating perturbation experiments to delineate the relevant causal circuit, as well as the application of other tools more readily available in rodents. In addition, advances in stimulus design and analysis have aided studying the relevant neural encoding. In complement to ongoing nonhuman primate studies, these newly available model systems and tools place the field at an exciting time that suggests that the dynamical circuit mechanisms underlying accumulation of evidence could soon be revealed.

Introduction

When we face a difficult decision, and are therefore uncertain as to what the best choice is, we are slow to make up our minds; but when faced with an easy decision, we are fast. This experience from daily life is one of the most common behavioral observations in decision-making, and applies in a remarkably wide array of different types of decisions, ranging from perceptual decisions [1], to numerical comparison decisions [2], to social decisions [3], to visual search decisions [4,5], to gambling decisions [6], to memory retrieval decisions [7], to lexical retrieval decisions [8], to social decisions [3], to value-based decisions [9–15]. A conceptually simple model, introduced many decades ago in the behavioral literature [7,16–19] has been able to account very well for the observation across all the above decision-making domains. As a result, this model, known as the "evidence accumulation" or "evidence integration" model, has become widely adopted as a succinct yet powerful behavioral-level description of core decision-making processes.

Publisher's Disclaimer: This is a PDF file of an unedited manuscript that has been accepted for publication. As a service to our customers we are providing this early version of the manuscript. The manuscript will undergo copyediting, typesetting, and review of the resulting proof before it is published in its final citable form. Please note that during the production process errors may be discovered which could affect the content, and all legal disclaimers that apply to the journal pertain.

The central idea of the model is that as a subject is forming a decision, evidence for or against different possible choices is gradually accumulating in the subject's mind; the final value of this accumulated evidence then drives the decision itself-- for example, committing to a particular choice by asking whether the accumulated evidence lies to one side or another of a reference value which we will label as the "decision boundary" (see Fig. 1a). In this model, when the evidence is strong, the accumulator's value quickly diverges away from the decision boundary, and it rapidly become easy to say on which side of the decision boundary it is. Whereas when the evidence is weak, the value of the evidence accumulator meanders away from the decision boundary only very gradually, leading to slower, more difficult decisions.

Here we will first briefly describe studies with non-human primates into the neural basis of the evidence accumulator. Our main focus will then be on more recent work using rodents, proposed as a complementary model system with which to unravel the mechanistic circuit dynamics underlying the phenomenon. We limit ourselves to two-alternative decisions (for multi-alternative decisions, see [20]), and will not address work in humans (see [21]).

Seminal studies in non-human primates

Starting in 1996, Michael Shadlen, William Newsome and colleagues (in addition to related parallel work from Jeffrey Schall's group [22]), began a series of highly influential electrophysiological experiments that suggested a connection between neural responses and evidence accumulation. Using monkeys trained to perform a visual perceptual decision-making task in which the experimenters could titrate each trial's difficulty ("random dot motion discrimination" (RDM) task; [23,24], Shadlen and Newsome found that during the process of decision formation, neurons in the lateral intraparietal (LIP) subregion of posterior parietal cortex (PPC) had firing rates that appeared to ramp up in time. Critically, when averaged across easy trials, firing rates ramped slowly; but when averaged across difficult trials, firing rates ramped much more sharply--precisely as expected of the evidence accumulator ([25–27]; Fig. 1b). This was the first time anyone had observed a signal *inside the brain* that matched what had been predicted for many years by the widespread accumulator model. Their seminal finding led to the proposal that there may be a 1-to-1 relationship between PPC firing rates and the value of the evidence accumulator.

Work in several laboratories (e.g., [28]) has uncovered similar firing rate patterns in multiple brain regions, most prominently in the frontal eye fields (FEF) [4,29–31], but also in other regions (dorsolateral prefrontal cortex, [29]; superior colliculus [32–34]; and striatum [35]; Fig. 1c). To date, causal perturbation studies of these areas with the primate RDM task have been limited, with only three existing published studies, all using only unilateral electrical microstimulation. Gold and Shadlen used microstimulation in the FEF to conclude that "developing oculomotor commands may reflect the formation of the monkey's direction judgement," but made no conclusions about the causal role of the FEF itself [36,37]. In the striatum, Ding and Gold found mixed effects, with contralateral responses speeded up even while, surprisingly, ipsilateral responses were favored [38]. Finally, in the PPC, Hanks et al. found that whole-trial microstimulation could be interpreted as adding a constant offset to the accumulator variable [39], although they acknowledged that their data "cannot rule out

the possibility that some of the observed effects are due to antidromic activation of other areas." This concern applies to all microstimulation studies, and is consistent with the later claim that microstimulation acts mainly through activation of axons [40] (but see [41]). The antidromic activation issue is not a concern with more recent inactivation studies to be described below. Despite these many studies, unraveling the neural circuit mechanisms underlying the accumulation process has remained a challenge, and major fundamental questions are still unanswered or even unaddressed. Which brain regions are necessary or not necessary for the evidence accumulation process? Where in the brain is evidence accumulation actually computed--perhaps in brain region(s) not yet recorded from? What is the macro- and micro- circuit architecture that supports the gradual accumulation computation?

New tools and animal models for studying gradual evidence accumulation

To assist in addressing these questions, Brunton et al. developed a rat model of a perceptual decision task dependent on gradual accumulation of evidence ("Poisson clicks" task, [42]; see also [43,44]). Rats were presented with two simultaneous trains of randomly timed auditory clicks, one to their left, the other to their right, and were trained to orient at the end of the stimulus towards the side that had the greater total number of pulses. As with the primate RDM task, trial difficulty could be titrated by the experimenters, in this case by controlling the left: right ratio of the Poisson rates that generated the random pulses. The highly variable yet very precisely known timing of the pulses ensured that stimulus space was thoroughly explored. It provided rich information that gave statistical power to a trial-by-trial decision process model that took into account the timing of each individual click, and that allowed estimating multiple parameters of the decision process. Together with model-free analyses that supported the main conclusions, the model provided strong behavioral evidence that the rats were indeed using gradual accumulation of evidence to perform the task. The door was open for using a cheap, small, tractable mammalian animal model for studies of evidence accumulation.

Using rats trained to perform the Poisson clicks task, Hanks, Kopec, et al. recorded from two cortical regions in the rat brain, the rat PPC and the frontal orienting fields (FOF; [45]). On the basis of their connectivity with other brain regions and their physiological properties, these had been suggested as analogous or perhaps even homologous to the two key primate regions PPC and FEF, respectively [46,47]. As in the primate, Hanks, Kopec, et al. found that neurons in the rat regions had trial-averaged ramping firing rates during the process of decision formation. Further, the slope of the ramp was steeper for stronger evidence-- again, just as had been observed in the primate brain regions. Also as in the primate, traces were qualitatively similar across the two recorded regions (Fig. 2a,b). The similarities across species suggested that rats and primates might be using similar circuits and algorithms to solve the task. Whether this is indeed the case remains an open question, of course. But it became clear that a fruitful comparison and intellectual exchange across model species was possible.

The pulsatile nature of the Poisson click stimulus, combined with a model-based estimate of each trial's evolution of the accumulator, allowed Hanks, Kopec, et al. to perform two

analyses that in previous primate work had required additional experiments. First, by computing the click-triggered average firing rate, it was possible to estimate the impact of a single evidence pulse on a neuron's firing rate. A perfect accumulator would respond with a permanent change in its value. Both PPC and FOF were found to respond to an evidence pulse with sustained changes in firing rate, albeit with a slow decay in the case of the FOF (see Fig. 2c,d and see [48] for the corresponding primate PPC experiment). Second, the Brunton et al. model produced a moment-by-moment estimate of the gradually evolving value of the evidence accumulator. This meant that at each point in time, in each trial, Hanks et al. had both a measure of a neuron's firing rate and an estimate of the value of the evidence accumulator. Collating these paired measures across trials allowed building "tuning curves," plots of average firing rates as a function of the variable of interest, which in this case is the value of the accumulator.

We note that this model-based approach combines firing rate measurements with knowledge of the full, detailed within-trial dynamics of the sensory stimulus as well as the animal's behavior. It is thus very different to a variety of interesting methods that have focused entirely on statistical analyses of neuronal firing [28,49,50]. Future work combining both types of approaches will be valuable.

The new tuning curve analysis tool revealed that, on average across the population, tuning curves in both the PPC and the FOF were relatively stable during the decision formation. Thus, although the value of the accumulator changes during the course of a trial, the mapping from accumulator value to firing rate does not. But while tuning curves in PPC were found to be smooth functions of the value of the accumulator --indicating that the graded value of the accumulator could be read out from PPC firing, as also found in the primate PPC [51]-- the tuning curves in the FOF were much more step-like, with one cluster of firing rates for accumulator values to one side of the decision boundary, and a different cluster of firing rates for values on the other side of the decision boundary (see Fig. 2e). Thus what is best read out from FOF firing is not the graded value of the accumulator, but rather which categorical decision report the subject should make if the trial ended at that point in time. This suggests that the FOF may be more strongly involved in categorizing the value of the accumulator to drive the final choice than in computing the graded accumulation itself. The tuning curve analysis thus allowed distinguishing two regions with qualitatively similar responses, and began the task of differentiating which, out of the various steps in decision formation, each region might be more closely associated with. More specifically, the analysis suggested a simple serial organization in our fully trained rats, in which the graded value of the accumulator is computed in the PPC, and then read out from there to be turned into a more categorical representation in the FOF; which in turn eventually drives the categorical motor acts with which the animal reports its decision. In this serial model, activity in both regions is critical for task performance, but activity in the FOF is required specifically near the time when the animal is triggered to report its decision (the "GO cue"), since that is when the overall decision process needs to read out on which side of the decision boundary the accumulator value lies. As we describe below, this serial model turned out to be wrong.

Perturbation studies

One key advantage of rodents is that computerizing and parallelizing behavioral training across subjects can produce a sufficient number of trained animals to make multiple perturbation experiments viable. By their very nature, perturbation experiments alter the brain, and therefore tend to be used very sparingly when individual animals are expensive in terms of cost, time, or effort involved in their training. For example, producing 50 trained subjects to troubleshoot and deploy a new perturbation technique can be prohibitively expensive with non-human primates. In contrast, it is straightforwardly practical with computerized, parallelized rodent training.

The tens of milliseconds time resolution afforded by optogenetic inactivation, in this case using the virally-delivered, light-activated chloride pump eNpHR3.0 [52] injected unilaterally into the FOF, allowed probing and confirming the hypothesis about the specific times when activity in the FOF would be required (see Fig 2f; [45]). Model-based analysis of performance impairments caused by unilateral, as well as bilateral, hours-long pharmacological silencing of the FOF [53] provided further converging evidence for the hypothesized role for the FOF: namely, that the FOF's main role in the task lies in categorizing the value of the accumulator, an operation that occurs *after* the graded evidence accumulation itself. Ongoing experiments are further probing the hypothesis through temporally-specific bilateral optogenetic inactivation of the FOF.

In contrast to expectations, however, inactivation of the PPC produced entirely negligible effects on performance of the Poisson Clicks task [53]. This was the case even while the same inactivations had a large effect on interspersed control "free choice" trials [53], and even though PPC inactivations have a substantial effect on a different auditory task involving parametric working memory [54]. These data suggest that despite its encoding of the graded value of the accumulator, the PPC may play little to no role in decisions driven by accumulation of evidence. Supporting this idea, preliminary data reproduces the PPC inactivation finding in visual evidence accumulation tasks in primates [55], inactivation of the primate pulvinar, which is the PPC-projecting region of the thalamus, has no impact on decision choices in a related visual task [56], inactivation of the primate PPC has no effect on visual primate memoryguided tasks [57–60], inactivation of mouse PPC has no effect on a somatosensory memory-guided task [61], and temporally-specific inactivation of mouse PPC in a visual memory-guided task has an effect only during the sensory stimulus period, not the short-term memory maintenance period [62]. While the PPC could, perhaps, play a role in gradual accumulation as part of a much larger redundant circuit in the Poisson Clicks task, it is notable that out of 6 different rat regions probed so far (medial prefrontal cortex and superior colliculus, (Hanks, Yartsev, et al. unpublished data); anterior dorsal striatum and auditory striatum (Yartsey, Hanks, et al., unpublished data); and FOF and PPC), the only region for which silencing has no impact on task behavior is the PPC. Consequently, we consider the PPC to have the lowest likelihood of being an important center of the causally relevant circuit for gradual evidence accumulation.

In contrast to primate visual tasks and the rodent auditory Poisson Clicks task, but consistent with Goard et al.'s preliminary optogenetic inactivation data [62], hours-long

pharmacological inactivation of the PPC does impair rodent visual short-term memory or evidence accumulation tasks [63,64]. Immediately posterior to the rodent PPC are a set of individually small visual areas, collectively referred to as secondary visual cortex (V2; [65–67]). With the precise definition of rodent PPC and the location of its border with these V2 areas still highly uncertain and a matter of active research [68–70], the observed effects from inactivations targeting rodent PPC could perhaps be due to inactivation spillover into one of the many small V2 areas. The appeal of this possibility is that it would make results across mouse, rat, and monkey consistent, for the difference in results between rodent visual tasks versus primate visual tasks or rodent non-visual tasks would be explained by the fact that spillover into the relevant sensory cortices is not a concern in the latter two cases. Even if this view were correct, the role of any such V2 sub-area in accumulation of evidence tasks remains to be determined and could not simply be basic sensory processing of visual stimuli [63]. Perhaps like the FOF, which we have posited as being required for reading the *output* of the accumulator, this region might be required for providing visual sensory *input* to the accumulator, even while not being involved in accumulation itself.

Further clarity on the role of the PPC will greatly benefit from better definitions and knowledge about the topography of brain regions near the location currently referred to as "rodent PPC." We advocate reserving the name "posterior parietal cortex (PPC)" in rodents to regions displaying physiological and connectivity profiles similar to those of the primate PPC, for example, no impact from inactivation on visual accumulation of evidence tasks [55], little anatomical input from primary visual cortex [71,72], and strong connectivity with frontal regions, as is the case with the anterior, as opposed to the posterior, zones currently referred to as "rat PPC" [70].

So where is the accumulator circuit? Preliminary hints and research outlook

To study the macro- and micro-circuitry causally responsible for gradual accumulation, we must first know which brain regions are part of that circuitry. A region that is part of the relevant causal circuit would be expected to simultaneously satisfy three initial criteria: (1) inactivation of the region should have an effect on task performance. More specifically, because there are different sequential steps involved in the decision formation and gradual accumulation is only one of them, (2) perturbation during temporally-specific periods that coincide with the evidence accumulation period should impact performance. And (3) the graded value of the accumulator should be encoded in that region's neural activity (as assessed, for example, through the tuning curve method of Hanks, Kopec, et al.). As described above, neither the FOF nor the PPC satisfy all three criteria. Where in the brain are the region(s) that do?

The data analytic and experimental tools developed to study the PPC and the FOF are now being applied to several rat brain regions, an endeavour facilitated by high-throughput rodent training. Current unpublished data (Yartsev, Hanks, et al.) suggests that while some subcortical regions may be like the FOF in playing an important role in the task, but not in gradual accumulation per se, at least one other subcortical region appears to satisfy all three

criteria, and may thus become the first identified node in the causal circuit for evidence accumulation. Assuming that this preliminary conjecture is correct, by following projections to, and projections from, the identified region, we may finally be able to delineate and establish the relevant causal circuit. Initial suggestions indicate that this may include brain regions never before examined in the context of accumulation of evidence.

We are thus in a particularly exciting time for research into the neural underpinnings of the evidence accumulator: once the relevant causal circuit is established, we will at long last be able to probe, and empirically distinguish, the many theoretical accounts that have been proposed [73–86]. In addition, methods to apply cellular-resolution imaging of neural activity to this problem, with which to powerfully examine multi-neuronal population codes [87], have recently been developed [88,89], and the many genetic tools available in mice may soon be brought to bear as well [90–92].

Finally, we emphasize our strong view that the developments in rodent model species described in this review, while very exciting, are complementary to efforts in non-human primates, in which sophisticated behavioral manipulations can be most powerfully employed to reveal important aspects of neural encoding and function, very fine control over both sensory and motor responses can be achieved in the visuomotor domain (e.g., [37,93]), and which are an important link towards relating neurophysiological findings to neural mechanisms in humans.

Acknowledgments

We gratefully acknowledge funding from NIH grant 1R01MH108358 and Simons Collaboration on the Global Brain award 325403 to CDB and NIH grant F32MH098572 to TDH.

References

- Gold JI, Shadlen MN. The neural basis of decision making. Annu. Rev. Neurosci. 2007; 30:535–574. [PubMed: 17600525]
- 2. Sigman M, Dehaene S. Parsing a cognitive task: a characterization of the mind's bottleneck. PLoS Biol. 2005; 3:e37. [PubMed: 15719056]
- 3. Krajbich I, Hare T, Bartling B, Morishima Y, Fehr E. A Common Mechanism Underlying Food Choice and Social Decisions. PLoS Comput. Biol. 2015; 11:e1004371. [PubMed: 26460812]
- 4. Purcell BA, Heitz RP, Cohen JY, Schall JD, Logan GD, Palmeri TJ. Neurally constrained modeling of perceptual decision making. Psychol. Rev. 2010; 117:1113–1143. [PubMed: 20822291]
- 5. Heitz RP, Schall JD. Neural mechanisms of speed-accuracy tradeoff. Neuron. 2012; 76:616–628. [PubMed: 23141072]
- Busemeyer, JR. Time pressure and stress in human judgment and and decision making. Springer; 1993. p. 181-193.
- 7. Ratcliff R. A Theory of Memory Retrieval. Psychol. Rev. 1978; 85:59-108.
- 8. Anders R, Riès S, van Maanen L, Alario F-X. Evidence accumulation as a model for lexical selection. Cogn. Psychol. 2015; 82:57–73. [PubMed: 26375509]
- 9. Basten U, Biele G, Heekeren HR, Fiebach CJ. How the brain integrates costs and benefits during decision making. Proc. Natl. Acad. Sci. U. S. A. 2010; 107:21767–21772. [PubMed: 21118983]
- 10. Milosavljevic M, Malmaud J, Huth A, Koch C, Rangel A. The Drift Diffusion Model can account for the accuracy and reaction time of value-based choices under high and low time pressure. Judgement and Decision Making. 2010; 5:437–449.

11. Krajbich I, Armel C, Rangel A. Visual fixations and the computation and comparison of value in simple choice. Nat. Neurosci. 2010; 13:1292–1298. [PubMed: 20835253]

- Cavanagh JF, Wiecki TV, Cohen MX, Figueroa CM, Samanta J, Sherman SJ, Frank MJ. Subthalamic nucleus stimulation reverses mediofrontal influence over decision threshold. Nat. Neurosci. 2011; 14:1462–1467. [PubMed: 21946325]
- 13. Hunt LT, Kolling N, Soltani A, Woolrich MW, Rushworth MFS, Behrens TEJ. Mechanisms underlying cortical activity during value-guided choice. Nat. Neurosci. 2012; 15:470–6. S1–3. [PubMed: 22231429]
- 14. Krajbich I, Lu D, Camerer C, Rangel A. The attentional drift-diffusion model extends to simple purchasing decisions. Front. Psychol. 2012; 3:193. [PubMed: 22707945]
- 15. Solway A, Diuk C, Córdova N, Yee D, Barto AG, Niv Y, Botvinick MM. Optimal behavioral hierarchy. PLoS Comput. Biol. 2014; 10:e1003779. [PubMed: 25122479]
- 16. Stone M. Models for Choice-Reaction Time. Psychometrika. 1960; 25:251–260.
- 17. Vickers D. Evidence for an accumulator model of psychophysical discrimination. Ergonomics. 1970; 13:37–58. [PubMed: 5416868]
- 18. Laming DRJ. Information theory of choice-reaction times [Internet]. 1968 [no volume].
- Ratcliff R, McKoon G. The diffusion decision model: theory and data for two-choice decision tasks. Neural Comput. 2008; 20:873–922. [PubMed: 18085991]
- Churchland AK, Ditterich J. New advances in understanding decisions among multiple alternatives. Curr. Opin. Neurobiol. 2012; 22:920–926. [PubMed: 22554881]
- Ratcliff, R., Smith, PL. Modeling simple decisions and applications using a diffusion model. In: Busemeyer, JR.Wang, Z.Townsend, JT., Eidels, A., editors. Oxford Handbook of Computational and Mathematical Psychology. Oxford University Press; 2015.
- Hanes DP, Schall JD. Neural control of voluntary movement initiation. Science. 1996; 274:427–430. [PubMed: 8832893]
- 23. Newsome WT, Paré EB. A selective impairment of motion perception following lesions of the middle temporal visual area (MT). J. Neurosci. 1988; 8:2201–2211. [PubMed: 3385495]
- 24. Newsome WT, Britten KH, Movshon JA. Neuronal correlates of a perceptual decision. Nature. 1989; 341:52–54. [PubMed: 2770878]
- Shadlen MN, Newsome WT. Motion perception: seeing and deciding. Proc. Natl. Acad. Sci. U. S. A. 1996; 93:628–633. [PubMed: 8570606]
- 26. Shadlen MN, Newsome WT. Neural basis of a perceptual decision in the parietal cortex (area LIP) of the rhesus monkey. J. Neurophysiol. 2001; 86:1916–1936. [PubMed: 11600651]
- 27. Roitman JD, Shadlen MN. Response of neurons in the lateral intraparietal area during a combined visual discrimination reaction time task. J. Neurosci. 2002; 22:9475–9489. [PubMed: 12417672]
- 28. Bollimunta A, Totten D, Ditterich J. Neural dynamics of choice: single-trial analysis of decision-related activity in parietal cortex. J. Neurosci. 2012; 32:12684–12701. [PubMed: 22972993]
- Kim JN, Shadlen MN. Neural correlates of a decision in the dorsolateral prefrontal cortex of the macaque. Nat. Neurosci. 1999; 2:176–185. [PubMed: 10195203]
- Ding L, Gold JI. Neural correlates of perceptual decision making before, during, and after decision commitment in monkey frontal eye field. Cereb. Cortex. 2012; 22:1052–1067. [PubMed: 21765183]
- **31. Mante V, Sussillo D, Shenoy KV, Newsome WT. Context-dependent computation by recurrent dynamics in prefrontal cortex. Nature. 2013; 503:78–84. The authors demonstrate the importance of population codes and their dynamics. They analyze responses of primate FEF neurons from monkeys performing an accumulation of evidence task in which the sensory feature providing evidence to be accumulated was different in different blocks of trials. As in many previous prefrontal cortex studies, the authors report a highly mixed, distributed population code. But, going beyond that, they also describe intriguing structure in the neural responses, by which different stimulus features control evidence accumulation in a goal-directed manner, within a line attractor representation. An artificial neural network trained to perform the same task shows qualitatively similar properties, and suggests that stimulus feature selection is controlled not by gating information, but by subtly adapting the direction of the dynamical flow of neural activity. [PubMed: 24201281]

32. Horwitz GD, Newsome WT. Separate signals for target selection and movement specification in the superior colliculus. Science. 1999; 284:1158–1161. [PubMed: 10325224]

- 33. Horwitz GD, Newsome WT. Target selection for saccadic eye movements: prelude activity in the superior colliculus during a direction-discrimination task. J. Neurophysiol. 2001; 86:2543–2558. [PubMed: 11698541]
- 34. Ratcliff R, Cherian A, Segraves M. A comparison of macaque behavior and superior colliculus neuronal activity to predictions from models of two-choice decisions. J. Neurophysiol. 2003; 90:1392–1407. [PubMed: 12761282]
- 35. Ding L, Gold JI. Caudate encodes multiple computations for perceptual decisions. J. Neurosci. 2010; 30:15747–15759. [PubMed: 21106814]
- 36. Gold JI, Shadlen MN. Representation of a perceptual decision in developing oculomotor commands. Nature. 2000; 404:390–394. [PubMed: 10746726]
- 37. Gold JI, Shadlen MN. The influence of behavioral context on the representation of a perceptual decision in developing oculomotor commands. J. Neurosci. 2003; 23:632–651. [PubMed: 12533623]
- 38. Ding L, Gold JI. Separate, causal roles of the caudate in saccadic choice and execution in a perceptual decision task. Neuron. 2012; 75:865–874. [PubMed: 22958826]
- 39. Hanks TD, Ditterich J, Shadlen MN. Microstimulation of macaque area LIP affects decision-making in a motion discrimination task. Nat. Neurosci. 2006; 9:682–689. [PubMed: 16604069]
- 40. Histed MH, Bonin V, Reid RC. Direct activation of sparse, distributed populations of cortical neurons by electrical microstimulation. Neuron. 2009; 63:508–522. [PubMed: 19709632]
- 41. Tehovnik EJ, Slocum WM. Two-photon imaging and the activation of cortical neurons. Neuroscience. 2013; 245:12–25. [PubMed: 23603308]
- **42. Brunton BW, Botvinick MM, Brody CD. Rats and humans can optimally accumulate evidence for decision-making. Science. 2013; 340:95–98. This paper exploited the statistical power of pulse-based tasks. Using sensory stimuli with randomly-chosen but precisely-known timing of pulses of evidence, the authors constructed a moment-by-moment model of the decision process. Fitting this model to behavioral data allowed simultaneously estimating multiple decision process parameters, including the first separate estimates of noise associated with incoming sensory evidence versus noise associated with holding a value in the accumulator. Surprisingly, the accumulator was found to be essentially lossless and noiseless, both for human data and for rat data. The model also provided a moment-by-moment estimate of the trajectory of the accumulator within each trial, a feature capitalized on in ref. [45]. And together with some model-free analyses, the model-based analysis provided strong behavioral evidence that, like primates, rodents can gradually accumulate evidence for decision-making. [PubMed: 23559254]
- 43. Sanders JI, Kepecs A. Choice ball: a response interface for two-choice psychometric discrimination in head-fixed mice. J. Neurophysiol. 2012; 108:3416–3423. [PubMed: 23019000]
- 44. Raposo D, Sheppard JP, Schrater PR, Churchland AK. Multisensory decision-making in rats and humans. J. Neurosci. 2012; 32:3726–3735. [PubMed: 22423093]
- **45. Hanks TD, Kopec CD, Brunton BW, Duan CA, Erlich JC, Brody CD. Distinct relationships of parietal and prefrontal cortices to evidence accumulation. Nature. 2015; 520:220–223. The authors used the moment-by-moment estimates of the accumulator's value provided by the model of ref. [42], together with neural recordings in rat PPC and FOF, to estimate "tuning curves" for mentally accumulated evidence, i.e., plots of firing rate versus value of the accumulator. They found important differences between the two regions. A closely related analysis, reaching qualitatively similar conclusions as in rat PPC, was performed for primate PPC in ref. [51]. (The corresponding experiment in primate FEF remains to be carried out.) Hanks, Kopec, et al. also transiently inactivated the FOF and found that only inactivations at the end of the accumulation period had an impact on performance, consistent with the proposal that the FOF's main role is to transform the final output of the graded accumulator into a categorical choice signal. [PubMed: 25600270]
- 46. Whitlock JR, Sutherland RJ, Witter MP, Moser M-B, Moser EI. Navigating from hippocampus to parietal cortex. Proc. Natl. Acad. Sci. U. S. A. 2008; 105:14755–14762. [PubMed: 18812502]
- 47. Erlich JC, Bialek M, Brody CD. A cortical substrate for memory-guided orienting in the rat. Neuron. 2011; 72:330–343. [PubMed: 22017991]

48. Huk AC, Shadlen MN. Neural activity in macaque parietal cortex reflects temporal integration of visual motion signals during perceptual decision making. J. Neurosci. 2005; 25:10420–10436. [PubMed: 16280581]

- 49. Churchland AK, Kiani R, Chaudhuri R, Wang X-JJ, Pouget A, Shadlen MN. Variance as a signature of neural computations during decision making. Neuron. 2011; 69:818–831. [PubMed: 21338889]
- *50. Latimer KW, Yates JL, Meister MLR, Huk AC, Pillow JW. Single-trial spike trains in parietal cortex reveal discrete steps during decision-making. Science. 2015; 349:184–187. When trial-averaged firing rates look like smooth ramps, is that because individual trials were smooth ramps, or because individual trials looked like sharp steps and the stepping times varied across trials? Focusing on the statistics of spike times, and using the computational power afforded by modern graphical processing units (GPUs), the authors were able to explore the myriad possible trajectories that noisy ramps or noisy stepping times could take. Their analysis suggested that about three-quarters of apparently smooth ramp-like firing rates in primate PPC are statistically better described as being composed of individual trials in which firing rates change in a sudden step. [PubMed: 26160947]
- **51. Kira S, Yang T, Shadlen MN. A neural implementation of Wald's sequential probability ratio test. Neuron. 2015; 85:861–873. The authors extended a previous study in which monkeys had learned to associate sequentially presented visual shapes with the likelihood of reward being either on the left or on the right [94]. Kira et al. took the Yang et al.'s task further in at least two critical directions: First, after being presented, each shape disappeared from the screen, forcing the subjects to accumulate the evidence from different shapes mentally. Since the experimenters knew the value of the evidence that each shape was associated with, and they knew the precise sequence of presented shapes on each trial, they had an estimate of the accumulated evidence at each timepoint in each trial, and could correlate that with simultaneously recorded firing rates of PPC to build tuning curves, as in ref. [45]. Kira et al. found a smooth, graded relationship between PPC neuron firing rates and accumulated evidence (as also found in rat PPC in ref. [45]). Second, the subjects were allowed to choose the timing of their response (a "reaction time" task). This time was found to be closely correlated to the time at which PPC neurons reached a fixed firing rate threshold. [PubMed: 25661183]
- 52. Gradinaru V, Zhang F, Ramakrishnan C, Mattis J, Prakash R, Diester I, Goshen I, Thompson KR, Deisseroth K. Molecular and cellular approaches for diversifying and extending optogenetics. Cell. 2010; 141:154–165. [PubMed: 20303157]
- *53. Erlich JC, Brunton BW, Duan CA, Hanks TD, Brody CD. Distinct effects of prefrontal and parietal cortex inactivations on an accumulation of evidence task in the rat [Internet]. Elife. 2015; 4 This manuscript reports the first pharmacological inactivations of brain regions within the decision-making circuit of Fig. 1c while subjects perform an accumulation of evidence task. Using a model-based analysis, they found that the impact of FOF inactivations was consistent with the idea that the FOF's main role is to transform the final output of the graded accumulator into a categorical choice signal. Surprisingly, however, both unilateral and bilateral inactivations of the PPC had minimal impact on performance of Brunton et al.'s Poisson Clicks task. Intriguingly, unilateral inactivation of the PPC did cause a noticeable ipsilateral bias when carried out simultaneously with bilateral FOF inactivation.
- Akrami A, El-Hady A, Kopec CD, Brody CD. Time dependent involvement of Posterior Parietal and Prefrontal cortex in a rat auditory parametric working memory task. Society for Neuroscience Abstracts. 2015; 624.09
- Yates J, Katz L, Park IM, Pillow JW, Huk A. Dissociated functional significance of choice-related activity across the primate dorsal stream. COSYNE abstract. 2014
- Komura Y, Nikkuni A, Hirashima N, Uetake T, Miyamoto A. Responses of pulvinar neurons reflect a subject's confidence in visual categorization. Nat. Neurosci. 2013; 16:749–755. [PubMed: 23666179]
- Chafee MV, Goldman-Rakic PS. Inactivation of parietal and prefrontal cortex reveals interdependence of neural activity during memory-guided saccades. J. Neurophysiol. 2000; 83:1550–1566. [PubMed: 10712479]
- 58. Wardak C, Olivier E, Duhamel J-R. Saccadic target selection deficits after lateral intraparietal area inactivation in monkeys. J. Neurosci. 2002; 22:9877–9884. [PubMed: 12427844]

 Liu Y, Yttri EA, Snyder LH. Intention and attention: different functional roles for LIPd and LIPv. Nat. Neurosci. 2010; 13:495–500. [PubMed: 20190746]

- Wilke M, Kagan I, Andersen RA. Functional imaging reveals rapid reorganization of cortical activity after parietal inactivation in monkeys. Proc. Natl. Acad. Sci. U. S. A. 2012; 109:8274– 8279. [PubMed: 22562793]
- *61. Guo ZV, Li N, Huber D, Ophir E, Gutnisky D, Ting JT, Feng G, Svoboda K. Flow of cortical activity underlying a tactile decision in mice. Neuron. 2014; 81:179–194. The authors developed and demonstrated the use of a powerful inactivation technique, in which they rendered transparent the skull of head-fixed, behaving mice. Using transgenic mice that expressed the light-activated excitatory opsin ChR2 in all GABAergic neurons then allowed them to locally inactivate (to ~1mm radius resolution) any dorsal brain site of their choosing by shining a laser through the transparent skull. This makes very high-throughput inactivation screening of different brain regions possible. The authors use their technique to show that, in a mouse oriented licking task guided by somatosensory short-term memory, barrel cortex is required during the sensory stimulus period but the frontal cortical region ALM is required during the memory retention period. [PubMed: 24361077]
- 62. Goard M, Pho G, Sur M. Optical dissection of cortical circuits underlying short-term memory. Society for Neuroscience Abstracts. 2014; 558.23
- 63. Harvey CD, Coen P, Tank DW. Choice-specific sequences in parietal cortex during a virtual-navigation decision task. Nature. 2012; 484:62–68. [PubMed: 22419153]
- 64. Raposo D, Kaufman MT, Churchland AK. A category-free neural population supports evolving demands during decision-making. Nat. Neurosci. 2014; 17:1784–1792. [PubMed: 25383902]
- 65. Coogan TA, Burkhalter A. Hierarchical organization of areas in rat visual cortex. J. Neurosci. 1993; 13:3749–3772. [PubMed: 7690066]
- 66. Montero VM. Retinotopy of cortical connections between the striate cortex and extrastriate visual areas in the rat. Exp. Brain Res. 1993; 94:1–15. [PubMed: 8335065]
- 67. Wang Q, Burkhalter A. Area map of mouse visual cortex. J. Comp. Neurol. 2007; 502:339–357. [PubMed: 17366604]
- 68. Reep RL, Chandler HC, King V, Corwin JV. Rat posterior parietal cortex: topography of corticocortical and thalamic connections. Exp. Brain Res. 1994; 100:67–84. [PubMed: 7813654]
- 69. Corwin JV, Reep RL. Rodent posterior parietal cortex as a component of a cortical network mediating directed spatial attention. Psychobiology. 1998; 26:87–102.
- Wilber AA, Clark BJ, Demecha AJ, Mesina L, Vos JM, McNaughton BL. Cortical connectivity maps reveal anatomically distinct areas in the parietal cortex of the rat. Front. Neural Circuits. 2014; 8:146. [PubMed: 25601828]
- 71. Felleman DJ, Van Essen DC. Distributed hierarchical processing in the primate cerebral cortex. Cereb. Cortex. 1991; 1:1–47. [PubMed: 1822724]
- 72. Markov NT, Ercsey-Ravasz MM, Ribeiro Gomes AR, Lamy C, Magrou L, Vezoli J, Misery P, Falchier A, Quilodran R, Gariel MA, et al. A weighted and directed interareal connectivity matrix for macaque cerebral cortex. Cereb. Cortex. 2014; 24:17–36. [PubMed: 23010748]
- 73. Seung HS, Lee DD, Reis BY, Tank DW. Stability of the memory of eye position in a recurrent network of conductance-based model neurons. Neuron. 2000; 26:259–271. [PubMed: 10798409]
- Usher M, McClelland JL. The time course of perceptual choice: the leaky, competing accumulator model. Psychol. Rev. 2001; 108:550–592. [PubMed: 11488378]
- 75. Wang X-J. Probabilistic decision making by slow reverberation in cortical circuits. Neuron. 2002; 36:955–968. [PubMed: 12467598]
- Koulakov AA, Raghavachari S, Kepecs A, Lisman JE. Model for a robust neural integrator. Nat. Neurosci. 2002; 5:775–782. [PubMed: 12134153]
- 77. Mazurek ME, Roitman JD, Ditterich J, Shadlen MN. A role for neural integrators in perceptual decision making. Cereb. Cortex. 2003; 13:1257–1269. [PubMed: 14576217]
- Goldman MS, Levine JH, Major G, Tank DW, Seung HS. Robust persistent neural activity in a model integrator with multiple hysteretic dendrites per neuron. Cereb. Cortex. 2003; 13:1185– 1195. [PubMed: 14576210]

79. Machens CK, Romo R, Brody CD. Flexible control of mutual inhibition: a neural model of two-interval discrimination. Science. 2005; 307:1121–1124. [PubMed: 15718474]

- 80. Fransén E, Tahvildari B, Egorov AV, Hasselmo ME, Alonso AA. Mechanism of graded persistent cellular activity of entorhinal cortex layer v neurons. Neuron. 2006; 49:735–746. [PubMed: 16504948]
- 81. Wong K-F, Wang X-J. A recurrent network mechanism of time integration in perceptual decisions. J. Neurosci. 2006; 26:1314–1328. [PubMed: 16436619]
- 82. Bogacz R, Brown E, Moehlis J, Holmes P, Cohen JD. The physics of optimal decision making: a formal analysis of models of performance in two-alternative forced-choice tasks. Psychol. Rev. 2006; 113:700–765. [PubMed: 17014301]
- 83. Bogacz R, Gurney K. The basal ganglia and cortex implement optimal decision making between alternative actions. Neural Comput. 2007; 19:442–477. [PubMed: 17206871]
- 84. Goldman MS. Memory without feedback in a neural network. Neuron. 2009; 61:621–634. [PubMed: 19249281]
- 85. Goldman MS, Compte A, Wang XJ. Neural integrator models. Encyclopedia of neuroscience. 2009; 6:165–178.
- Lim S, Goldman MS. Balanced cortical microcircuitry for maintaining information in working memory. Nat. Neurosci. 2013; 16:1306–1314. [PubMed: 23955560]
- 87. Kiani R, Cueva CJ, Reppas JB, Newsome WT. Dynamics of neural population responses in prefrontal cortex indicate changes of mind on single trials. Curr. Biol. 2014; 24:1542–1547. [PubMed: 24954050]
- 88. Scott BB, Brody CD, Tank DW. Cellular resolution functional imaging in behaving rats using voluntary head restraint. Neuron. 2013; 80:371–384. [PubMed: 24055015]
- 89. Scott BB, Constantinople CM, Erlich JC, Tank DW, Brody CD. Sources of noise during accumulation of evidence in unrestrained and voluntarily head-restrained rats [Internet]. Elife. 2015; doi: 10.7554/eLife.11308
- Odomene K, Brown AM, Kaufman MT, Chruchland AK. Disrupting inhibition in posterior parietal cortex reduces decision accuracy. Society for Neuroscience Abstracts. 2014; 358.21
- Harvey CD. Neuronal microcircuit dynamics in the mouse parietal cortex. Society for Neuroscience Abstracts. 2015; 458.07
- 92. Koay SA, Deverett B, Thiberge SY, Brody CD, Tank DW. A mouse navigation and accumulation of visual evidence task. Society for Neuroscience Abstracts. 2015; 442.15
- 93. Bennur S, Gold JI. Distinct representations of a perceptual decision and the associated oculomotor plan in the monkey lateral intraparietal area. J. Neurosci. 2011; 31:913–921. [PubMed: 21248116]
- 94. Yang T, Shadlen MN. Probabilistic reasoning by neurons. Nature. 2007; 447:1075–1080. [PubMed: 17546027]

Brody and Hanks

Highlights

 Gradual accumulation of evidence, core to decision-making, can be studied in rodents

- Pulse-based sensory evidence stimuli facilitate behavioral and neural analysis
- High-time-resolution inactivation helps identify role of different brain regions
- Differential role of the five brain regions studied so far is being distinguished
- Causal circuit for accumulation involves only some of the regions studied so




Figure 1. Evidence accumulator models and associated circuits. a, Schematic of evidence accumulation process, here illustrated for a case when the subject must decide between orienting Left or Right. As the decision process unfolds, noisy evidence favoring one choice (RIGHT) adds to the accumulator while evidence favoring the other choice (LEFT) subtracts from the accumulator. The sign of the accumulated evidence when the subject is asked to report their decision dictates the resulting decision choice. Trials with strong evidence that more consistently favors one choice over the other result in steeper slopes on average, and the accumulator will soon be far away from the decision boundary, so easy decisions can be made quickly. Weaker, less consistent evidence will result in meandering trajectories with shallower slopes on average, and even after lengthy accumulation periods, the accumulator may not be far from the decision boundary, leading to slow, more error-prone decisions. In tasks in which the subject determines the duration of the decision process, known as "reaction time tasks," the subject commits to a decision when the evidence reaches a bound (+C or -C in the figure); the reaction time is determined by when the bound is reached, and the decision choice is given by which bound was reached. b, Average neural responses from monkey PPC (area LIP) during the period of decision formation in the random dot motion discrimination task [27]. After a delay, responses exhibit ramping response profiles with

slopes that depend on stimulus strength. Stronger motion leads to sharper slopes and weaker motion to shallower slopes. This corresponds to the average trends predicted by the evidence accumulator model. **c**, Diagram of interconnected brain regions that have been demonstrated to exhibit responses profiles correlated with accumulating evidence. These areas thus serve as candidates to be involved in the evidence accumulation process.




Figure 2.
Characteristics of rat PPC and FOF during accumulation of evidence based decision making, from [45]. **a**, Average neural responses in rat PPC during the Poisson Clicks task. Trials are grouped by average stimulus strength. Similar to monkey PPC, responses exhibit ramping profiles that depend on stimulus strength. **b**, Same as **a** for FOF. **c**, Click-triggered average responses for rat PPC during the Poisson clicks task. Individual clicks have a measurable and sustained influence on responses in PPC. **d**, Same as **c** for FOF. Individual clicks also produced a sustained response, with a magnitude that slowly but significantly decayed over hundreds of milliseconds. **e**, Time-average population comparison of tuning curves for

accumulating evidence in PPC and FOF. PPC shows a smoothly graded relationship, while FOF shows a sharper dependence on the sign of the accumulator value. **f**, Bias caused by 500-ms unilateral inactivation of FOF with halorhodopsin during one of four epochs of the task: before the stimulus (red), during the first half of the stimulus (yellow), during the second half of the stimulus (green), or during the movement period (blue). Only peri-choice perturbation of FOF has a significant effect on decision making. A further experiment using half the inactivation time period (250 ms) reached the same conclusion [45].